

Project 1: Text Retrieval

Information retrieval is the process of searching and returning relevant documents for a query from a collection. This project focuses on a text retrieval task, and includes two subtasks:

Task 1: Document Full Ranking

In this task, you are expected to rank documents based on their relevance to the question, where documents can be retrieved from the full document collection provided. It models a scenario where you are building an end-to-end retrieval system.

Task 2: top-k re-ranking

We provide you with an initial ranking of documents per question, and you are expected to re-rank the documents in terms of their relevance to the question. This is a very common real-world scenario since many end-to-end systems are implemented as retrieval, followed by top-k re-ranking. The re-ranking subtask allows you to focus on re-ranking only without needing to implement an end-to-end system.

Dataset

We provide you a corpus of 1,471,405 documents with the following information:

Split	# of examples
Task 1 train	532,751
Task 2 train	10
Task 1 test	6980
Task 2 test	33

In the dataset folder you find the following files:

- `"corpus.jsonl"`: This file contains a collection of documents we want to retrieve from. Each document has a unique id (`"_id"`), and its content (`"text"`).
- `"queries.jsonl"`: This file contains a collection of queries that we want to retrieve relevant documents for. Each query has a unique id (`"_id"`), and its content (`"text"`).

- `task1_train.tsv`: This file is the “train” split of the data for Task 1 containing relevant documents for a set of queries and their relevance scores (always 1).
- `task2_train.tsv`: This file is the “train” split of the data for Task 2 containing relevant documents for a set of queries and their relevance scores (between 0 to 4).
- `task1_test.tsv`: This file is the “test” split of the data for Task 1 containing the queries you need to retrieve documents for.
- `task2_test.tsv`: This file is the “test” split of the data for Task 2 containing a set of queries and their retrieved candidates where you need to rank them based on their relevance.

Project Evaluation

- 40% : Results - Metrics
- 30% : Code
 - Working code (20%)
 - Code quality and documentation (10%)
- 30% : 2-page Report
 - Originality of approach (10%)
 - Interpretation of results (10%)
 - Report presentation & clarity (10%)

Regarding the Results part, we evaluate the performance of your implementation using two criteria:

1. Computation time (T): One evaluation criterion is the time it takes for your code to retrieve documents for the given test set in Task 1 and rerank the retrieved documents in Task. If this time is less than a **T0 = 10mins**, you get the full points of this criteria, and for computation time more than T0, you get penalized using this formula **$\max(0, 1 - T0/T)$** . This criterion counts for **1/4** of your grade for the Results part.
2. Weighted average of retrieval and reranking scores (S):

Task 1: For each query given in the Task 1 test set (`task1_test.tsv`), you have to retrieve 10 relevant documents from the given corpus. We evaluate the performance of your retrieval system using **Recall@10** metric.

Task 2: For each query given in the Task 2 test set (`task2_test.tsv`), you have to re-rank the retrieved documents in terms of their relevance to the query.

We evaluate the performance of your re-ranking system using **NDCG**¹ metric. The relevance scores can be any **real (non-negative) values**, and you are not limited to output values between 0 and 3. The metric is not sensitive to the exact values, and it cares about the rankings.

The final evaluation score is computed as a weighted average of the scores computed for these two tasks:

$$S = w_1 * recall@10 + w_2 * NDCG$$

where $w_1 = 0.8$ and $w_2 = 0.2$.

This criterion counts for **3/4** of your grade for the Results part. From this portion, the last 10% of the teams in the Kaggle competition get 4/6 of the grade, the top 10% get 6/6, and the groups in the middle get 5/6.

We ran a simple baseline with a final score of **0.47** to give you a better idea of where your approach stands. **You must achieve this performance at least to be considered in the grading explained before.**

Requirements:

- You can use any supervised or unsupervised methods in the project. However, for the supervised methods, you have to train the model yourself (using the given labeled data), and you are not allowed to use any trained models publicly available.
- You are not allowed to use TF-IDF implementation from sklearn or any other libraries.

Submission

Data distribution, running codes, and submission are all done using the Kaggle platform. Please create a Kaggle account with your **EPFL email**. Similar to colab, you can run notebooks (with or without GPU) on Kaggle for investigating data, as well as the final **submission** notebook. This notebook should write the submission file as one of the outputs. We will give more details about the final deliverables during the Thursday exercise session.

¹ Normalized Discounted Cumulative Gain (NDCG)

In order to join the competition (with your **EPFL email Kaggle**), please use this [invitation link](#).

Note: Throughout the competition, please don't make your notebooks public, and keep them **private** throughout the competition!